

## Spring 2020 - CS 477/577 - Introduction to Computer Vision

# Assignment Eight

**Due: 11:59pm (\*) Thursday, April 9.**

(\*) There is grace until 8am the next morning, as the instructor will not grade assignment before then. However, once the instructor starts grading assignments, no more assignments will be accepted.

**Weight: Approximately 5 points**

**This assignment can be done in groups of 2-3. You should create one PDF for the entire group, but everyone should hand in a copy of the PDF to help me keep track. Make it clear on the first page of your PDF who your group is. Group reports are expected to be higher quality than individual ones.**

---

### General instructions

You can use any language you like for this assignment, but unless you feel strongly about it, you might consider continuing with Matlab.

You need to create a PDF document that tells the story of the assignment, copying into it output, code snippets, and images that are displayed when the program runs. Even if the question does not remind you to put the resulting image into the PDF, if it is flagged with (\$), you should do so. I should not need to run the program to verify that you attempted the question. See

<http://kobus.ca/teaching/assignment-instructions.pdf>

for more details about doing a good write-up. While it takes work, it is well worth getting better and more efficient at this.

**30% of the assignment grade is reserved for exposition.**

### Learning goals

- Understand practical image filtering
- Understand basics of smoothing and edge detection
- Understand combining consecutive filters and separable filters
- Learn about linking edge points into edges, and more specifically about non-maximal suppression, following normal-to-gradient direction, and hysteresis (grads, choice one)
- Learn more about using convolution to implement correlation to look for known patterns (templates) (grads, choice two)

### Assignment specification

This assignment has four parts, one which is a few things to try in Matlab to get started but has no deliverables, one of which is required for undergrads, and two for which graduate students should choose one. If you do both, you will be eligible for modest extra credit.

## Part A (Recommended, but do not write it up!)

1) In Matlab (or an alternative, but then function names might be different) use `conv()` and `flip()` to do the 1D examples in the notes; 2) Further investigate commutativity using small random kernels and signals; 3) Investigate the question in class about whether we always get the flipped version of the output from convolution when we use correlation instead of convolution (as suggested from the one example in lecture 17, slide 9) by trying a few cases with small random kernels and signals; and (for fun) 4) Matlab “help conv” (paraphrased) says that if A and B store the coefficients of polynomials, then `conv(A,B)` is the coefficients of the terms in their product. Think for a moment why that is the case, why the flipping is key here, and see if you get the right answer for  $(1+2x+3x^2)*(3-3x)$ .

## Part B (Required for both undergrads and grads)

To keep things simple, we will focus on black and white (gray level) images. This means that an image is basically a matrix. Thus, both in Matlab and in C you might find it easiest to get the image into a matrix, work on the matrix, and then convert the matrix back to an image for display.

The input image for this assignment is <http://kobus.ca/teaching/cs477/data/climber.tiff> (also in D2L Data)

1. Write a program to show the magnitude of the gradient of the input image using convolution to implement finite differences. You will want to scale the brightness of your result so that it looks reasonable. Your program should display your result, which should also be put into your writeup (\$).

You can use builtin / library routine(s) for convolution in Matlab / C. Or you may prefer to implement convolution yourself to help make sure you really understand it. (A naive implementation in Matlab will run much slower than the built in function--you should understand why). However, you should **not** use the Matlab routine `imgradient`, as it obscures what is going on a bit too much for pedagogy.

2. Now find a threshold for the gradient magnitude to determine real edges. Make a new image which is white if the gradient magnitude exceeds that threshold, and black (zero) otherwise. Likely you will find that there does not exist a perfect threshold which identifies exactly all points on all significant edges. Don't spend too much time tweaking the threshold, but do find a reasonable value. Put the result into your writeup (\$).
3. Now use convolution to smooth the image with a Gaussian mask with sigma set at 2 pixels. First create the mask using relatively low level operations (either loops and the Gaussian formula, or perhaps using functions like `meshgrid` and `mvnpdf`). The point is to make sure you understand filter creation, including judging a reasonable size and scaling it so that it sums to one. Provide a surface plot of your filter (\$). Then smooth your image using a convolution routine which does not assume anything about the mask (e.g., `conv2`). In particular, do **not** use `imgaussfilt` as it hides too much of what is going on. Put the result into your writeup (\$).
4. Now we will do blurring followed by edge detection. Apply the procedure from (1) and (2) to find edges in the blurred image that you computed in (3). Thus, for the image computed in (3), do edge detection as in (1) and find a threshold as in (2) and output a binary image based on that threshold (as in (2)). Put the result into your writeup (\$).
5. Now we will combine blurring and edge detection into one filter. In particular, convolve the image with an appropriate single mask to implement **both** smoothing by a sigma set to 4 pixels and finite difference operation in the X direction. This filter is the composition of two filters. Do

the same for the Y direction. Combine the two results, and put the new gradient image into your writeup (§). Verify for yourself (i.e., this is simply a check---don't hand anything in) that this gives roughly the same result as doing the convolutions separately with the appropriate masks, which is like doing (4) with a sigma of 4 instead of 2.

6. A function of  $x$  and  $y$ ,  $f(x,y)$  is separable (in  $x$  and  $y$ ), if  $f(x,y)=g(x)h(y)$  for some  $g(x)$  and  $h(y)$ . If this is the case, then convolution with  $f(x,y)$  can be implemented as a 1D convolution by  $g(x)$  followed by a 1D convolution by  $h(y)$ . Algebraically specify  $g(x)$  and  $h(y)$  in the case of Gaussian blurring, and show that  $f(x,y)=g(x)h(y)$  (§). Re-implement the smoothing in (3) using this approach. Verify for yourself that this gives roughly the same result as before.

*Matlab users: This is just a matter of calling `conv2` with different parameters.*

Do you expect any speed difference? Why? Do you see any speed difference? (You may want to put the convolution in a loop, or use a bigger sigma and image to test for time differences). You might consider a different separable function other than the Gaussian smoother in case Matlab is too smart about your input. Put your description of what you found into your writeup (§).

## (Required for grad students only).

Choose between C1 and C2 below. If you do both of them, you will be eligible for modest extra credit.

### Part C1 (edge detection)

Implement the gradient based edge detector described in class. You will need to implement the maximum suppression part, as well as edge linking. Tell the TA the story of your attempt at this, how your edge detector performs with various choices for thresholds, sigmas, etc. (provide images with informative captions). Your linked edges should be colored in some way to demonstrate that sequence of edge points aggregated into edges is represented in your program. §

### Part C2 (OCR)

The following file has some rendered text (also in D2L Data)

[http://kobus.ca/teaching/cs477/data/lorem\\_ipsum.tiff](http://kobus.ca/teaching/cs477/data/lorem_ipsum.tiff)

Create templates for  $N$  letters by cutting small images containing just that letter out of the image. By default, consider  $N=8$ , but you can choose a different number if it seems more interesting to you or better in some way. Recall that it is usually a good idea to add a constant to such a filter so that the response to uniform image sections is zero (i.e., the filter sums to zero). Use correlation to find areas of high filter response to each of the filters. You will need to set a threshold for when the likelihood of misidentification is too great. Letters that do not meet the threshold should be classified as “other”. Also, if you have a hit at a given threshold for two letters, you will naturally choose the one with highest response. If you do not have any such conflicts, perhaps add a few letters to keep it interesting. You will want to visualize the results of your OCR system (§). In particular, you want to make it clear which letters are found, and what they are identified as, perhaps with some color-coded boxes overlaid on the original images. Again, you will have made some choices (e.g., a threshold), and the TA would love to know what you found, partly through nice pictures. Finally, provide a confusion matrix for your  $N$  letters and the extra category “other” (§).

*In a confusion matrix, the rows and columns both correspond to the items you are trying to classify (e.g., letters). So, if you have 8 letters and “other”, you will have 9 rows and 9 columns. The order of correspondence should be the same. Each element of the matrix counts how many*

*times the row letter was classified as a column letter. So, a perfect classifier only has non-zero elements on the diagonal. If we scale the matrix properly, we have an estimate of how likely a letter (say “Q” will be called any of the possible letters, and tells us about how “confused” the classifier is. This is very good for understanding in detail the kinds of mistakes your classifier is making.*

---

## **What to Hand In**

As usual, the main deliverable will be PDF document that tells the story of your assignment as described above. Ideally the grader can focus on that document, simply checking that the code exists, and seems up to the task of producing the figures and results in the document. So, you need hand in your code as well.